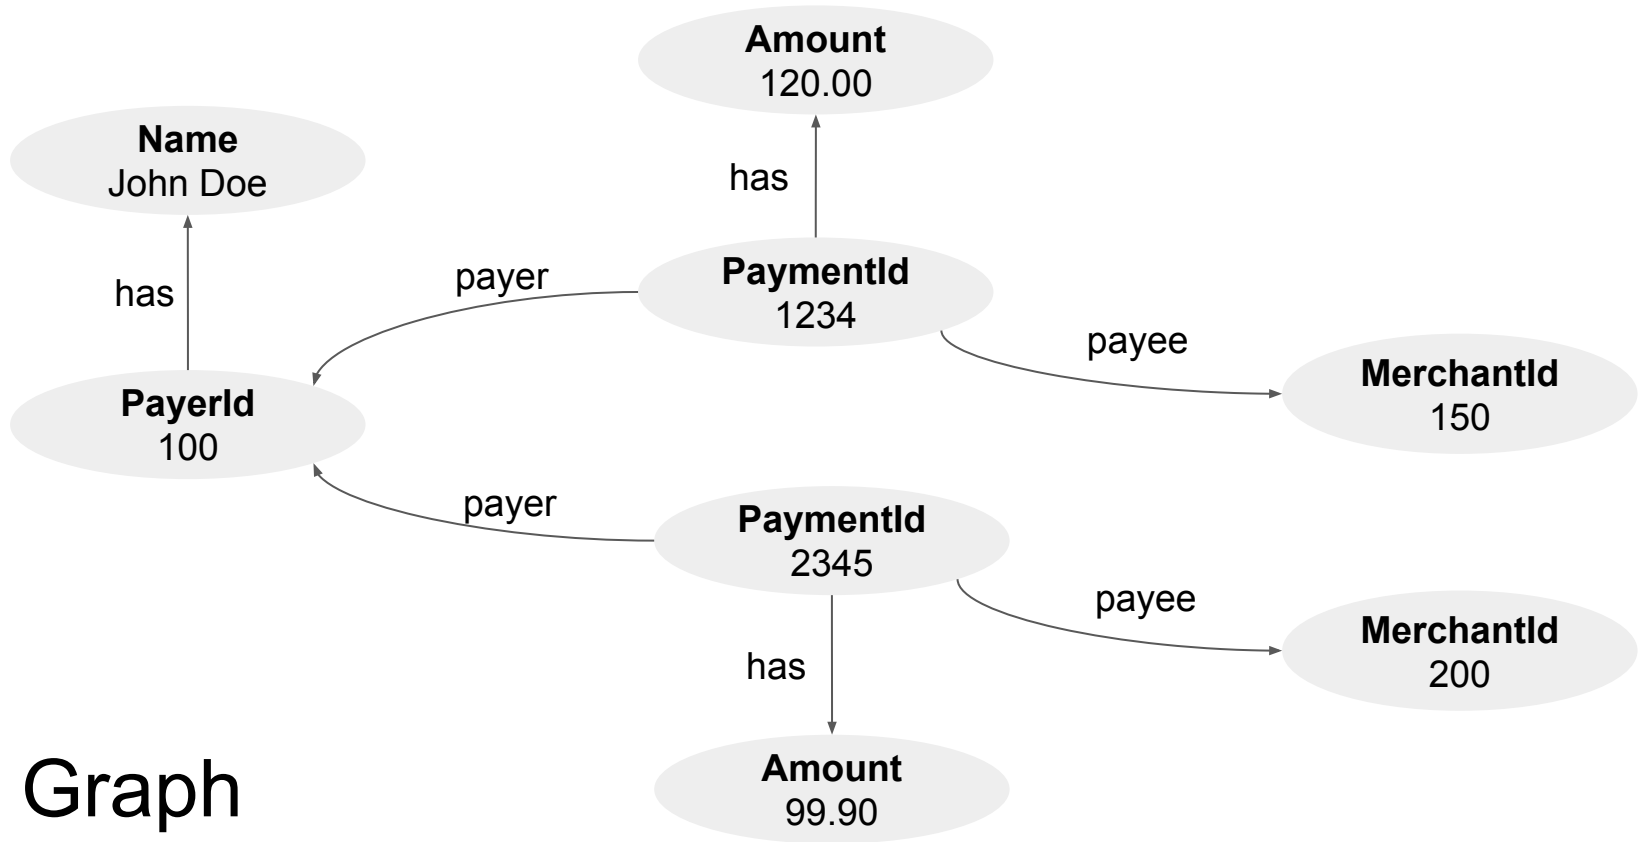


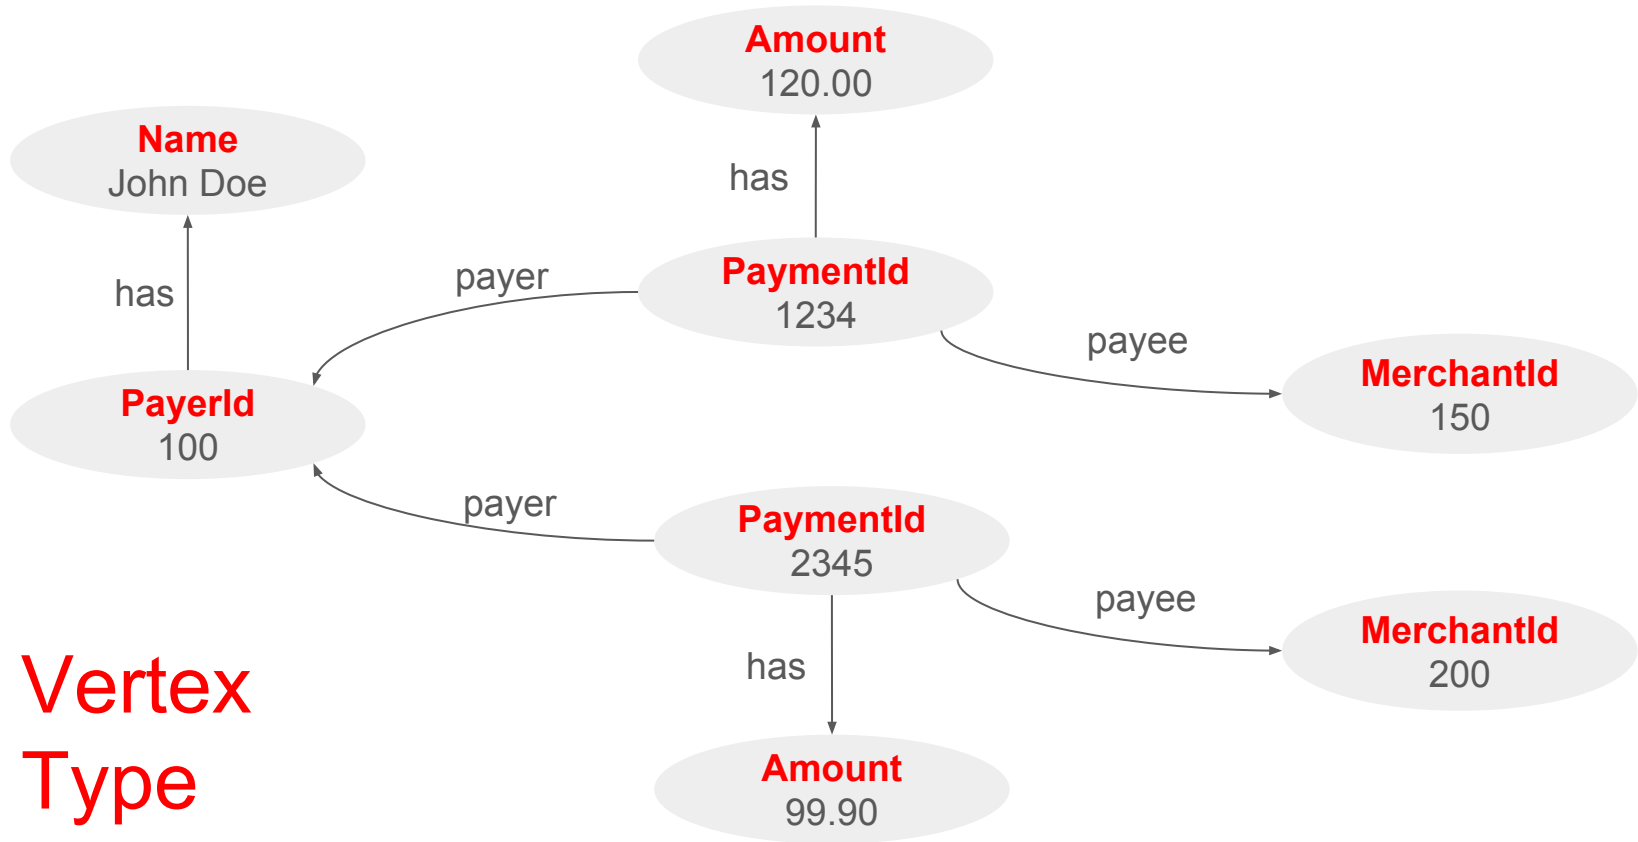
Cadenza

Cadenza

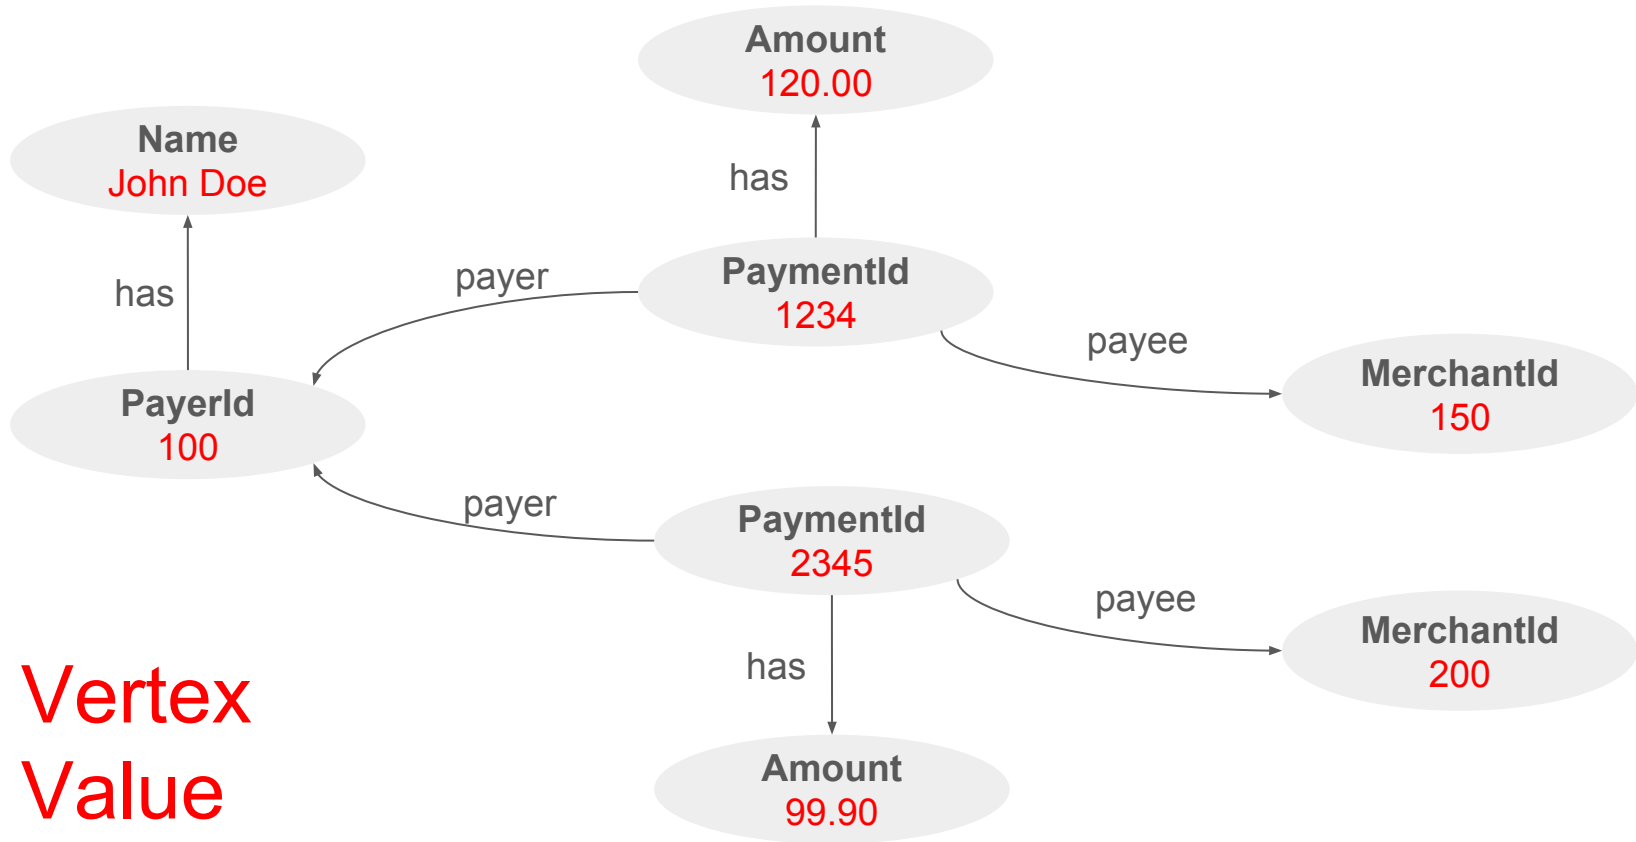
- (almost) in-memory database
 - Minimum JVM heap usage
 - Memory mapped file
 - Compact representation on file
 - Minimum record management overheads (minimum use of pointers, length fields, index structures, ...)
- Distributed (partitioned)
- Updatable
 - Chronological segmentation
 - Segment merge

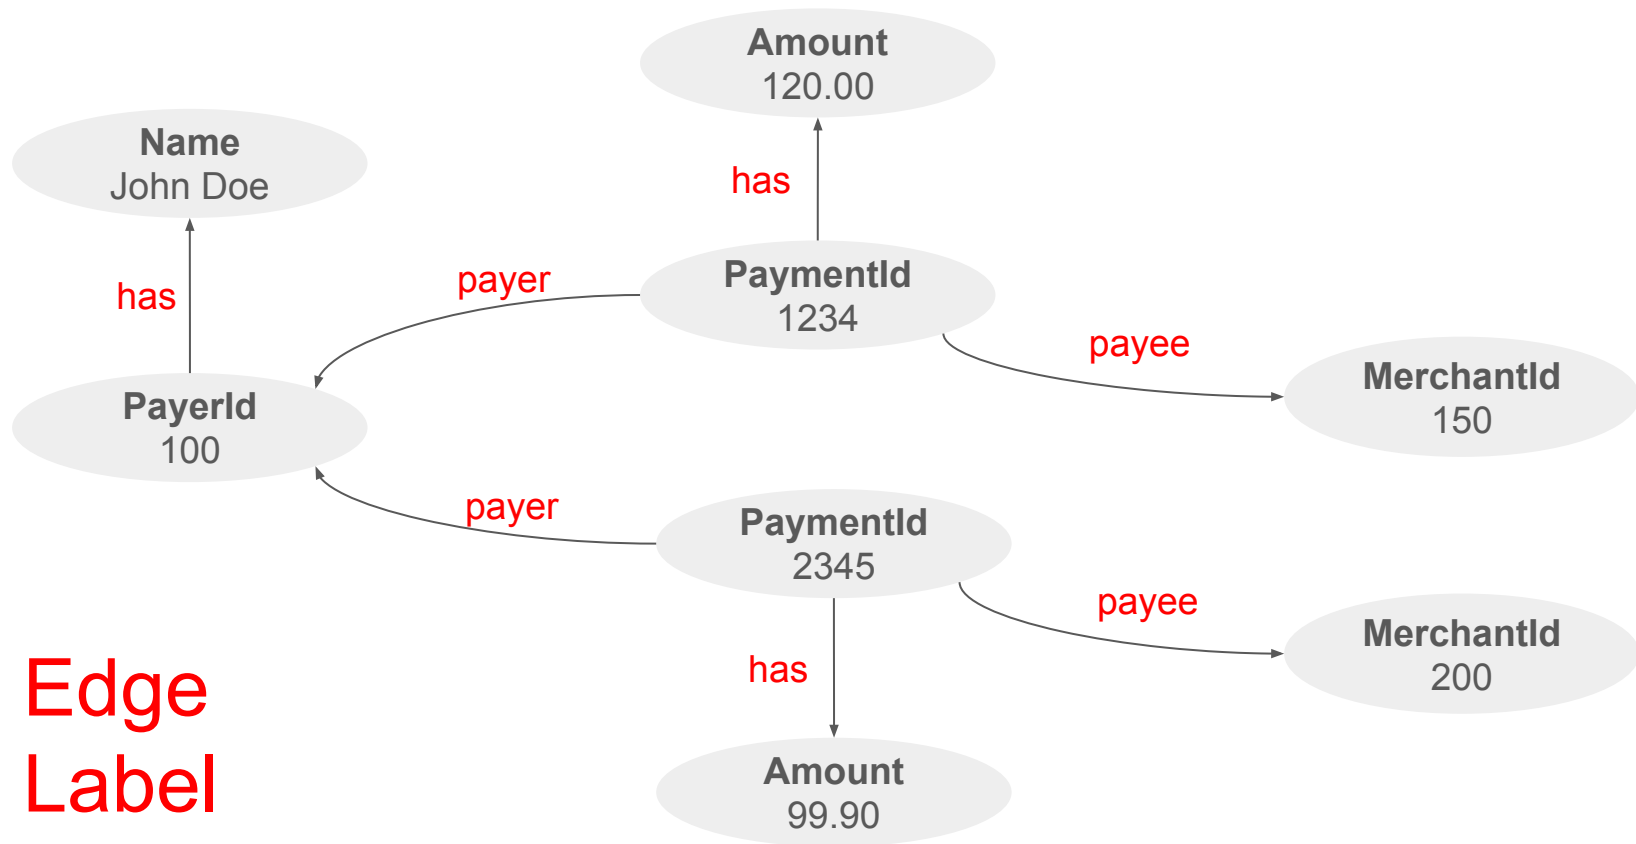


Graph



Vertex
Type





Edge
Label

Data Model

Vertex

- Typed: Vertex Type is a subtype of Long, Double, or String
- No attribute

Edge

- Labeled (named)
 - A label is independent from vertex types of source/destination
- Directed
- No attribute

In the example graph...

Vertex Type

- PayerId : Long
- MerchantId: Long
- Amount : Double
- Name : String

Edge Label (Edge Type)

- payer, payee, has

Query Language (WIP)

```
SELECT PID, AMNT, M
FROM {
    P Payer,
    M Merchant,
    PID PaymentId,
    AMNT Amount

    P <-[payer]- PID -[payee]-> M,
    PID -[has]-> AMNT
}
WHERE P = 100 AND AMNT > 100.00;
```

```
SELECT PID, AMNT, M
```

```
FROM {
```

```
    P Payer,
```

```
    M Merchant,
```

```
    PID PaymentId,
```

```
    AMNT Amount
```

```
    P <-[payer]- PID -[payee]-> M,
```

```
    PID -[has]-> AMNT
```

```
}
```

```
WHERE P = 100 AND AMNT > 100.00;
```

Subgraph Specification

```
SELECT PID, AMNT, M  
FROM {
```

Vertex Variable Declarations

```
    P Payer,  
    M Merchant,  
    PID PaymentId,  
    AMNT Amount
```

```
    P <-[payer]- PID -[payee]-> M,  
    PID -[has]-> AMNT  
}  
WHERE P = 100 AND AMNT > 100.00;
```

```
SELECT PID, AMNT, M
FROM {
```

**Vertex
Variables**

```
P Payer,
M Merchant,
PID PaymentId,
AMNT Amount
```

Vertex Variable Declarations

```
P <-[payer]- PID -[payee]-> M,
PID -[has]-> AMNT
}
WHERE P = 100 AND AMNT > 100.00;
```

```
SELECT PID, AMNT, M
FROM {
```

Vertex Variable Declarations

```
    P Payer,
```

```
    M Merchant,
```

```
    PID PaymentId,
```

```
    AMNT Amount
```

Vertex
Types

```
    P <-[payer]- PID -[payee]-> M,
```

```
    PID -[has]-> AMNT
```

```
}
```

```
WHERE P = 100 AND AMNT > 100.00;
```

```
SELECT PID, AMNT, M
FROM {
    P Payer,
    M Merchant,
    PID PaymentId,
    AMNT Amount
```

Subgraph Patterns

```
P <-[payer]- PID -[payee]-> M,
PID -[has]-> AMNT
```

```
}
WHERE P = 100 AND AMNT > 100.00;
```

Edge Pattern

Single Edge Pattern

srcVertexVar -[*edgeLabel*]-> *destVertexVar*

or

destVertexVar <-[*edgeLabel*]- *srcVertexVar*

Edge patterns can be chained

v1 -[**e1**]-> **v2** -[**e2**]-> **v3** <-[**e3**]- **v4**

More Operators are planned

- GROUP BY
- HAVING
- ORDER BY
- Aggregate functions (COUNT, SUM, AVG, MIN, MAX)

Subgraph Specification \approx Joined Table Expression

Edge Pattern \approx Two-Column Table


$$V1 \text{ --}[e]\text{--> } V2$$
$$\approx (V1, V2)$$

Edge Patterns are joined over shared Vertex Variables

$$V1 \text{ --}[e]\text{--> } V2, V2 \text{ --}[e]\text{--> } V3$$
$$\approx (V1, V2) \text{ join } (V2, V3) \text{ over } V2$$


Subgraph Spec \rightarrow Join Tree

$S \rightarrow T \rightarrow U \rightarrow V$

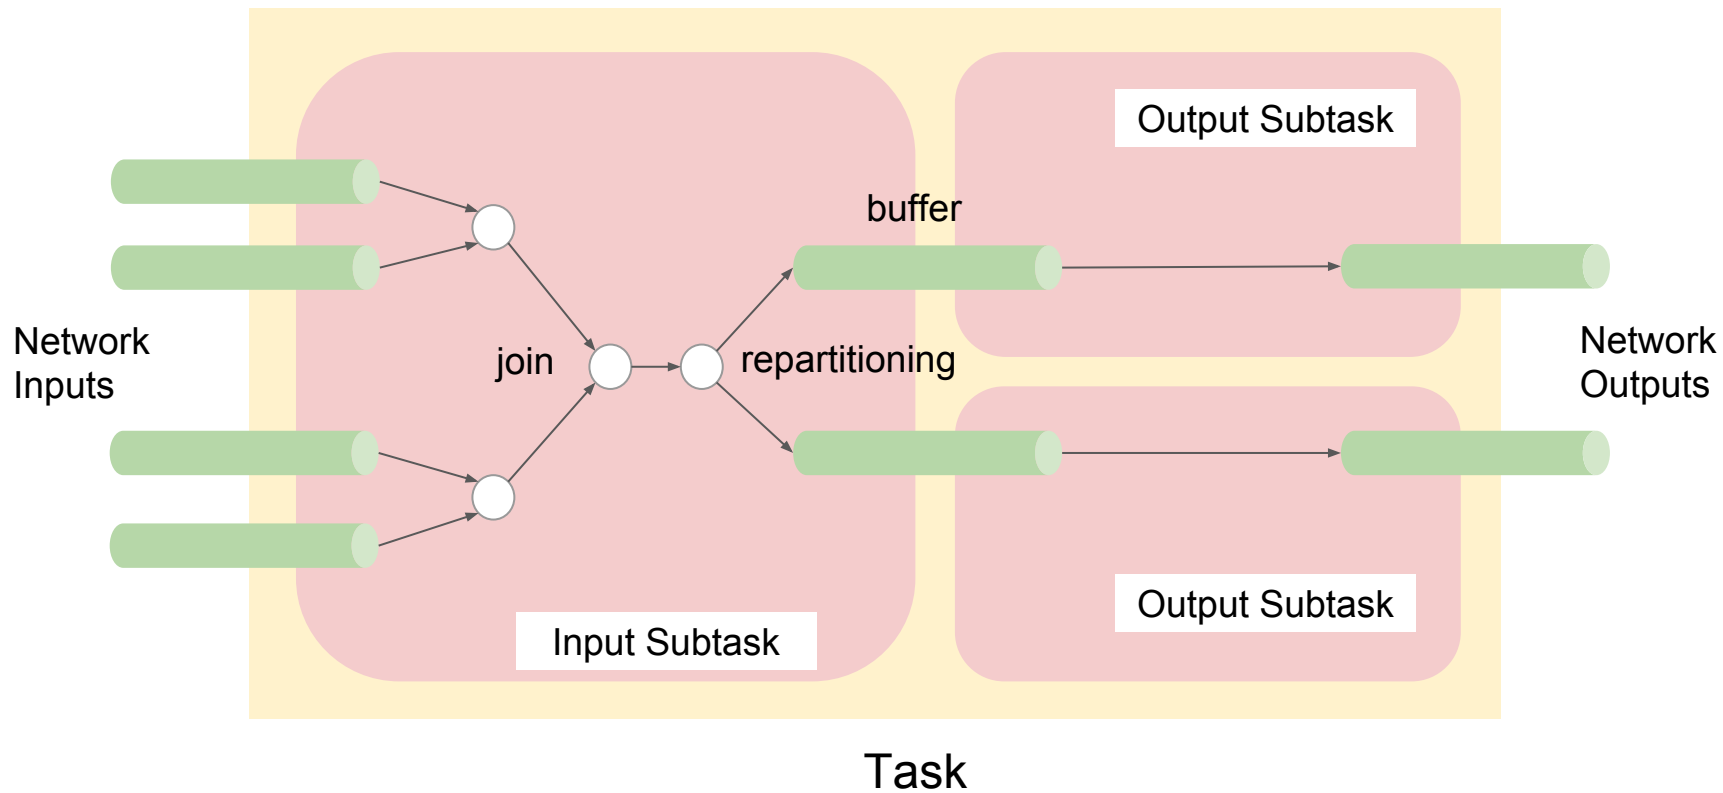


Distributed Processing

$S \rightarrow T \rightarrow U \rightarrow V$



Task Structure



Subtasks and Inputs/Outputs

a subtask \approx a thread

- Subtasks run in a thread pool
- A subtask should never block.

Input/Output may be suspended when data are not ready

- If input/output is not ready, a subtask is suspend.
- The suspended subtask is resumed when input/output becomes ready.

(see LocalProcessor)

Query Optimization

- Predicate Push-down
- Cost Based
- Exhaustive Search

Query Optimization

- Predicate Push-down
- Cost Based → Ad hoc, Naive Cost Model
- Exhaustive Search → Slow

Task Generation and Distribution

- A query plan is generated at a server
- A row source tree is generated from the plan
- TaskDefs are created from the row source tree
- A task is a combination of a TaskDef and a partition
- Tasks are assigned to servers
- Send tasks and bind variables to assigned servers

Storage

Edge List (source, destination, edge type)

- Compressed by removing repeating source vertices
 - Vertex List (list of <source vertex>)
 - Out-Edge List (list of <destination vertex, edge type>)
 - A succinct bit vector connects a vertex list and an out-edge list.
- Vertex Records and Out-Edge records are fixed size
- Vertex Records are sorted
- Out-Edge Records are grouped by source vertex, and sorted within the group

Inverted Index

- Indexed by destination vertices

Storage (cont.)

Inverted Index

- Indexed by hash values of destination vertices

Reverse Edge

- Reverse Edge is automatically created
 - An edge, $S \xrightarrow{e} D$, creates a reverse edge $D \xrightarrow{e^{-1}} S$ implicitly.
 - A reverse edge is use by the engine internally.

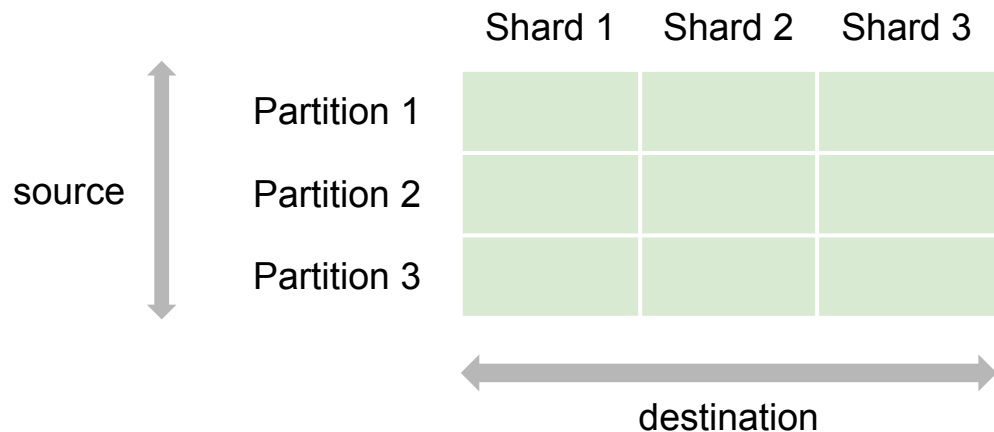
String Pool

String Vertex Value is a variable length data.

- Create a list of string values (sorted)
- Store pointer to a string in the pool in a vertex record.

Partitions and Shards

- Hash-partitioned by source vertices
- A partition is hash-partitioned into shards by destination vertices



Segments

A shard is chronologically segmented.

- Two in-memory segments
 - Writable in-memory segment (new data is written to it)
 - Read-only in-memory segment (it will be dumped to the disk)
- On-disk segments
 - A segment contains a Vertex list, an Out-edge list, a string pool, an inverted index
 - Segments are merged when there are too many.
 - We try to balance the merged segment sizes. (see MergePolicyImpl)

TODOs

- Language Parser [WIP]
- Language Features
 - Aggregate (group by, having, aggregate functions)
 - Arithmetic operators
 - Additional primitive types (Datetime?)
 - Edge label variable?
- Data Dictionary (schema management)
- Client Library
 - Native networking vs gRPC
- Data Ingestion
 - Waltz integration?
- Cluster management (Config?, ZK integration?, Gossip Protocol?)

com.wepay.cadENZA.common

- message package
 - Network messages
- network package
 - Networking layer (on Netty, copied from Waltz)
- pipe package
 - Network pipe (rows flow between remote tasks through this)
 - Local pipe (rows flow between local tasks through this)
- task package
 - Task, Subtask, TaskProcessor
- type package
 - Vertex, Accessor, Bind variable, Value, ...

com.wepay.cadENZA.server.query

- engine package
 - RowDef (describes a row)
 - RowSources (data access, join, filter, sort) and RowInputs
 - RuntimeContextImpl (access to bind vars, routing table, storage, and network pipe)
 - Task implementation
- plan package
 - Query execution plan generation
 - Row source tree generation
- Query classes
 - Subgraph, Where, OrderBy, Select
- Query Processor class
 - LocalProcessor

com.wepay.cadENZA.server.storage

- segment package
 - In-memory segment
 - On-disk segment
- util package
 - ImmutableBST (immutable randomized binary search tree)
 - SuccinctBitVector
- Input classes
 - VertexInput, OutEdgeInput, EdgeInput, ...
- Partition class
- Shard class